

# 01\_components\_and\_visualization

June 28, 2025

## 1 Components and Visualization of Time Series

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
import os
import plotly.io as pio
pio.templates.default = "plotly_white"
import pandas as pd
from pathlib import Path
from tqdm.autonotebook import tqdm
from IPython.display import display, HTML
# %load_ext autoreload
# %autoreload 2
np.random.seed()
tqdm.pandas()
```

/tmp/ipykernel\_7717/1114571209.py:9: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See

[https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html)

```
from tqdm.autonotebook import tqdm
```

Components of a time series: 1. Trend 2. Seasonal 3. Cyclical 4. Irregular

$$Y = Trend + Seasonal + Cyclical + Irregular$$

$$Y = Trend * Seasonal * Cyclical * Irregular$$

### 1.1 Trend Component

- The trend is a long-term change in the mean of a time series.
- It is the smooth and steady movement of a time series in a particular direction.

### 1.2 Seasonal Component

- The seasonal component is the recurring pattern of a time series that repeats itself every period.
- It is the recurring ups and downs within a period.

### 1.3 Cyclical Component

- The cyclical component is the recurring pattern of a time series that repeats itself every period.
- It is the recurring ups and downs within a period.

```
[3]: DATA_FOLDER_PATH = "/home/bilal326/Time_Series/data/london_smart_meters"

block_df = pd.read_parquet(f"{DATA_FOLDER_PATH}/preprocessed/
↳london_smart_meters_merged_block_0-7.parquet")

[6]: #Converting to expanded form
from data_utils import compact_to_expanded

exp_block_df = compact_to_expanded(block_df[block_df.file=="block_7"],
↳timeseries_col = 'energy_consumption',\
                                static_cols = ["frequency", "series_length",\
↳stdorToU", "Acorn", "Acorn_grouped",\
                                "file"],time_varying_cols =
↳['holidays', 'visibility', 'windBearing', 'temperature','dewPoint',\
↳'pressure', 'apparentTemperature', 'windSpeed', 'precipType', 'icon',\
↳'humidity', 'summary'],\

↳    ts_identifier = "LCLid")

# Taking a single time series from the block
ts_df = exp_block_df[exp_block_df.LCLid=="MAC000193"].set_index("timestamp")

ts_df["weekday_name"] = ts_df.index.day_name()
ts_df["weekday"] = ts_df.index.weekday
ts_df["week"] = ts_df.index.isocalendar().week
ts_df["day"] = ts_df.index.day
ts_df["hour"] = ts_df.index.hour
ts_df["date"] = ts_df.index.date
ts_df["month"] = ts_df.index.month
ts_df["month_name"] = ts_df.index.month_name()
ts_df["year"] = ts_df.index.year

#Making ordered categoricals to make for sorted plots
ts_df['month_name'] = pd.Categorical(ts_df['month_name'],
↳categories=["January", "February", "March", "April", "May", "June", "July",
↳"August", "September", "October", "November", "December"], ordered=True)
ts_df['weekday_name'] = pd.Categorical(ts_df['weekday_name'],
↳categories=["Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
↳"Saturday", "Sunday"], ordered=True)
```

0%| | 0/50 [00:00<?, ?it/s]

100%| | 50/50 [00:00<00:00, 74.93it/s]

```
[7]: # !wget https://raw.githubusercontent.com/PacktPublishing/
      ↪Modern-Time-Series-Forecasting-with-Python/refs/heads/main/src/imputation/
      ↪interpolation.py
```

--2025-06-28 18:18:05--

https://raw.githubusercontent.com/PacktPublishing/Modern-Time-Series-  
Forecasting-with-Python/refs/heads/main/src/imputation/interpolation.py

Resolving raw.githubusercontent.com (raw.githubusercontent.com)...

185.199.109.133, 185.199.110.133, 185.199.111.133, ...

Connecting to raw.githubusercontent.com

(raw.githubusercontent.com)|185.199.109.133|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 6820 (6.7K) [text/plain]

Saving to: 'interpolation.py.1'

interpolation.py.1 100%[=====>] 6.66K --.-KB/s in 0.001s

2025-06-28 18:18:06 (5.10 MB/s) - 'interpolation.py.1' saved [6820/6820]

## 1. Line Charts

```
[8]: #Interpolating Missing values
      # from src.imputation.interpolation import SeasonalInterpolation

      ts_df['energy_consumption'] = ts_df['energy_consumption'].
      ↪interpolate(method='linear')
```

```
[9]: fig = px.line(ts_df, y="energy_consumption", title="Energy Consumption for_
      ↪MAC000193")
      # fig = format_plot(fig, ylabel="Energy Consumption")
      fig.show()
```

## 2. Rolling Average Plot (rolling avg)

When you have a long time series with high variation, as we have, the line plot can get a bit chaotic. One of the options to get a macro view of the time series in terms of trends and movement is to plot a smoothed version of the time series.

```
[10]: ts_df["rolling_monthly_avg"] = ts_df["energy_consumption"].
      ↪rolling(window=48*30).mean()

      fig = px.line(ts_df, y="rolling_monthly_avg", title="Rolling Monthly Average_
      ↪Energy Consumption for MAC000193")
      fig.show()
```

Another use for the line chart is to visualize two or more time series together and investigate any correlations between them.

```
[11]: import plotly.graph_objects as go

fig = go.Figure()

# Add first y-axis (Energy Consumption)
fig.add_trace(go.Scatter(
    x=ts_df.index,
    y=ts_df['energy_consumption'],
    name='Energy Consumption',
    yaxis='y1',
    line=dict(color='blue')
))

# Add second y-axis (Temperature)
fig.add_trace(go.Scatter(
    x=ts_df.index,
    y=ts_df['temperature'],
    name='Temperature',
    yaxis='y2',
    line=dict(color='red', dash='dash')
))

# Set layout with dual y-axes
fig.update_layout(
    title='Temperature and Energy Consumption',
    xaxis=dict(title='Time'),
    yaxis=dict(
        title=dict(text='Energy Consumption', font=dict(color='blue')),
        tickfont=dict(color='blue')
    ),
    yaxis2=dict(
        title=dict(text='Temperature', font=dict(color='red')),
        tickfont=dict(color='red'),
        anchor='x',
        overlaying='y',
        side='right'
    ),
    legend=dict(x=0.01, y=0.99),
    width=900,
    height=500
)

fig.show()
```

```
[12]: import plotly.graph_objects as go

# Create rolling averages (e.g., 48*30 for ~monthly)
ts_df['energy_rolling'] = ts_df['energy_consumption'].rolling(window=48*30).
    ↪mean()
ts_df['temp_rolling'] = ts_df['temperature'].rolling(window=48*30).mean()

fig = go.Figure()

# Rolling Energy Consumption (y-axis 1)
fig.add_trace(go.Scatter(
    x=ts_df.index,
    y=ts_df['energy_rolling'],
    name='Energy Consumption (Rolling Avg)',
    yaxis='y1',
    line=dict(color='blue')
))

# Rolling Temperature (y-axis 2)
fig.add_trace(go.Scatter(
    x=ts_df.index,
    y=ts_df['temp_rolling'],
    name='Temperature (Rolling Avg)',
    yaxis='y2',
    line=dict(color='red', dash='dash')
))

# Layout with dual axes
fig.update_layout(
    title='Rolling Monthly Averages: Temperature and Energy Consumption',
    xaxis=dict(title='Time'),
    yaxis=dict(
        title=dict(text='Energy Consumption (Rolling Avg)',
    ↪font=dict(color='blue')),
        tickfont=dict(color='blue')
    ),
    yaxis2=dict(
        title=dict(text='Temperature (Rolling Avg)', font=dict(color='red')),
        tickfont=dict(color='red'),
        anchor='x',
        overlaying='y',
        side='right'
    ),
    legend=dict(x=0.01, y=0.99),
    width=900,
    height=500
)
```

```
fig.show()
```

### Seasonal plots

A seasonal plot is very similar to a line plot, but the key difference here is that the x-axis denotes the “seasons,” the y-axis denotes the time series value

And the different seasonal cycles are represented in different colors or line types.

```
[13]: plot_df = ts_df[~ts_df.year.isin([2011, 2014])].groupby(["year",  
    ↪ "month_name"])[['energy_consumption', "temperature"]].mean().reset_index()  
  
fig = px.line(plot_df, x="month_name", y='energy_consumption', color="year",  
    ↪ line_dash="year", title="Seasonal Plot - Monthly")  
# fig = format_plot(fig, ylabel="Energy Consumption", xlabel="Month")  
fig.show()
```

/tmp/ipykernel\_7717/3793814189.py:1: FutureWarning:

The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.


```
[14]: import matplotlib.pyplot as plt  
import pandas as pd  
  
# Make sure datetime index  
ts_df['month'] = ts_df.index.month  
ts_df['hour'] = ts_df.index.hour  
  
# Monthly averages  
monthly_avg = ts_df.groupby('month')[['energy_consumption', 'temperature']].  
    ↪ mean()  
  
# Plot  
fig, ax1 = plt.subplots(figsize=(10, 5))  
  
color = 'tab:blue'  
ax1.set_xlabel('Month')  
ax1.set_ylabel('Energy Consumption', color=color)  
ax1.plot(monthly_avg.index, monthly_avg['energy_consumption'], color=color,  
    ↪ label='Energy Consumption')  
ax1.tick_params(axis='y', labelcolor=color)  
  
# Second y-axis
```

```

ax2 = ax1.twinx()
color = 'tab:red'
ax2.set_ylabel('Temperature', color=color)
ax2.plot(monthly_avg.index, monthly_avg['temperature'], color=color,
        linestyle='--', label='Temperature')
ax2.tick_params(axis='y', labelcolor=color)

plt.title('Seasonal Trend by Month: Energy Consumption and Temperature')
fig.tight_layout()
plt.show()

```



```

[15]: # Hourly averages
hourly_avg = ts_df.groupby('hour')[['energy_consumption', 'temperature']].mean()


fig, ax1 = plt.subplots(figsize=(10, 5))

color = 'tab:blue'
ax1.set_xlabel('Hour of Day')
ax1.set_ylabel('Energy Consumption', color=color)
ax1.plot(hourly_avg.index, hourly_avg['energy_consumption'], color=color)
ax1.tick_params(axis='y', labelcolor=color)

ax2 = ax1.twinx()
color = 'tab:red'
ax2.set_ylabel('Temperature', color=color)
ax2.plot(hourly_avg.index, hourly_avg['temperature'], color=color,
        linestyle='--')
ax2.tick_params(axis='y', labelcolor=color)

```

```
plt.title('Seasonal Trend by Hour of Day: Energy Consumption and Temperature')
fig.tight_layout()
plt.show()
```



```
[16]: import matplotlib.pyplot as plt
import pandas as pd

# Make sure datetime index is set
ts_df['month'] = ts_df.index.month
ts_df['year'] = ts_df.index.year

# Filter for 2012 and 2013
season_df = ts_df[ts_df['year'].isin([2012, 2013])]

# Monthly averages
monthly_avg = season_df.groupby(['year', 'month'])[['energy_consumption',
    ↪ 'temperature']].mean().reset_index()

# Pivot to get separate columns for plotting
energy_pivot = monthly_avg.pivot(index='month', columns='year',
    ↪ values='energy_consumption')
temp_pivot = monthly_avg.pivot(index='month', columns='year',
    ↪ values='temperature')

# Plot
fig, ax1 = plt.subplots(figsize=(10, 6))
```



```

# Left y-axis (Energy)
color_energy = 'tab:blue'
ax1.set_xlabel('Month')
ax1.set_ylabel('Energy Consumption', color=color_energy)
ax1.plot(energy_pivot.index, energy_pivot[2012], label='Energy 2012',
        color=color_energy)
ax1.plot(energy_pivot.index, energy_pivot[2013], label='Energy 2013',
        color=color_energy, linestyle='--')
ax1.tick_params(axis='y', labelcolor=color_energy)

# Right y-axis (Temperature)
ax2 = ax1.twinx()
color_temp = 'tab:red'
ax2.set_ylabel('Temperature', color=color_temp)
ax2.plot(temp_pivot.index, temp_pivot[2012], label='Temp 2012',
        color=color_temp)
ax2.plot(temp_pivot.index, temp_pivot[2013], label='Temp 2013',
        color=color_temp, linestyle='--')
ax2.tick_params(axis='y', labelcolor=color_temp)


# Title and legend
plt.title('Seasonal Monthly Trends: Energy Consumption and Temperature (2012 vs
        2013)')

# Combine legends from both axes
lines_1, labels_1 = ax1.get_legend_handles_labels()
lines_2, labels_2 = ax2.get_legend_handles_labels()
ax1.legend(lines_1 + lines_2, labels_1 + labels_2, loc='upper left')

# Month labels
ax1.set_xticks(range(1,13))
ax1.set_xticklabels(['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
                    'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])

plt.grid(True)
fig.tight_layout()
plt.show()

```



But when there are too many seasonal cycles to be plotted, it increases visual clutter. An alternative to a seasonal plot is a seasonal box plot.

```
[17]: plot_df = ts_df.groupby(["date", "weekday_name", "hour"])["energy_consumption"].
      ↪mean().reset_index().dropna()

fig = px.box(plot_df, y="energy_consumption", x="hour", log_y=True, title="Box_
      ↪Plot: Day of Month-Hourly Average")
fig.update_layout(height=700)
fig.show()
```

/tmp/ipykernel\_7717/1998492874.py:1: FutureWarning:

The default of `observed=False` is deprecated and will be changed to `True` in a future version of pandas. Pass `observed=False` to retain current behavior or `observed=True` to adopt the future default and silence this warning.

- the seasonal box plot is much more informative.
- The horizontal line in the box tells us about the median,
- the box is the interquartile range (IQR), and the points that are marked are the outliers.
- By looking at the medians, we can see that the peak consumption occurs from 9 A.M. onward.
- But the variability is also higher from 9 A.M.
- But the variability is also higher from 9 A.M.

```
[23]: import plotly.express as px

# Group by date, weekday name, and hour, then take the mean
plot_df2 = ts_df.groupby(["date", "weekday_name",
    ↪ "hour"])[ "energy_consumption"].mean().reset_index().dropna()

# Create pivot table: weekday as rows, hour as columns
pivot_table = plot_df2.pivot_table(index="weekday_name", columns="hour",
    ↪ values="energy_consumption", aggfunc='mean')

# Optional: order weekdays correctly
weekday_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
    ↪ 'Saturday', 'Sunday']
pivot_table = pivot_table.reindex(weekday_order)

# Plot heatmap
fig = px.imshow(
    pivot_table,
    labels=dict(x="Hour of Day", y="Day of Week", color="Energy Consumption"),
    aspect="auto",
    color_continuous_scale="bluered",
    title="Heatmap: Hourly Energy Consumption by Day of Week"
)

fig.update_layout(xaxis_title="Hour", yaxis_title="Weekday")
fig.show()
```

/tmp/ipykernel\_7717/3335086085.py:4: FutureWarning:

The default of `observed=False` is deprecated and will be changed to `True` in a future version of pandas. Pass `observed=False` to retain current behavior or `observed=True` to adopt the future default and silence this warning.

/tmp/ipykernel\_7717/3335086085.py:7: FutureWarning:

The default value of `observed=False` is deprecated and will change to `observed=True` in a future version of pandas. Specify `observed=False` to silence this warning and retain the current behavior

## 1.4 Autocorrelation plot

- Correlation means the strength of the relationship between two variables.
- Autocorrelation is the correlation between a time series and itself at different lags.
- Autocorrelation measures how a time series value at time  $t$  is **related to** its own **past values** (e.g.,  $t - 1, t - 2, \dots$ ).
- It tells you whether past values can help **predict** future ones and **how** they're related.

**Positive Autocorrelation Definition:** Positive autocorrelation means that if a value is **high**, the next value (at some lag  $k$ ) is also **likely to be high**, and if it's low, the next value is likely to be low too.

**Interpretation:**

- The series has a **momentum**-like behavior.
- Patterns **persist** — high follows high, low follows low.

**Example:** Imagine daily temperature. If today is warm, tomorrow is also likely to be warm.

**Visual:**

High  $\rightarrow$  High  $\rightarrow$  High  $\rightarrow$  ... Low  $\rightarrow$  Low  $\rightarrow$  Low  $\rightarrow$  ...

**Negative Autocorrelation Definition:** Negative autocorrelation means that if a value is **high**, the next value is more likely to be **low**, and vice versa.

**Interpretation:**

- The series shows **alternating** or **mean-reverting** behavior.
- A high value is followed by a dip and then a rise again.

**Example:** Think of a bouncing ball: every bounce goes up and then down.

**Visual:** High  $\rightarrow$  Low  $\rightarrow$  High  $\rightarrow$  Low  $\rightarrow$  ...

```
[22]: plot_df.head()
```

```
[22]:
```

	date	weekday_name	hour	energy_consumption
144	2012-01-01	Sunday	0	0.3770
145	2012-01-01	Sunday	1	0.0955
146	2012-01-01	Sunday	2	0.0380
147	2012-01-01	Sunday	3	0.0285
148	2012-01-01	Sunday	4	0.0360

### 1.4.1 Partial Autocorrelation

#### Partial Autocorrelation (PACF)

**Definition:** **Partial autocorrelation** measures the **direct effect** of a time series' past value at lag  $k$  on its current value, **after removing the effects of all shorter lags** (1 through  $k - 1$ ).

In simpler terms:

It tells you how much **extra predictive power** lag  $k$  adds **after accounting for** lags 1 to  $k - 1$ .

**Why is this useful?** In time series (especially AR models), just looking at autocorrelation (ACF) can be misleading — you might see high correlation at lag 3, but it might just be because lag 3 is indirectly related via lag 1 and 2.

Partial autocorrelation tells you:

- Does lag  $k$  have a **direct influence** on the current value?
- Or is its effect just due to its correlation with earlier lags?

**Example: Autocorrelation vs Partial Autocorrelation** Let's say we have a time series  $X_t$ , and:

- Lag 1 affects  $X_t$  directly.
- Lag 2 affects  $X_t$  **only because** it's correlated with Lag 1.

Then:

- ACF might show **high** correlation at both lag 1 and 2.
- **PACF** will show **high** only at lag 1, and **near-zero** at lag 2.

**In Practice: How it's used** In **AR (AutoRegressive)** models:

- **PACF plot** helps you choose the order  $p$  (how many lags to use).
- You look for the **last significant lag** in the PACF plot → that's often your AR model order.

### Quick Summary

Term	What it Measures	Includes Indirect Effects?	Used For
<b>ACF</b> (Autocorrelation Function)	Correlation of current value with past lags	Yes	Identifying MA (moving average) order
<b>PACF</b> (Partial Autocorrelation Function)	Direct correlation of lag $k$ , removing lags 1 to $k - 1$	No	Identifying AR (autoregressive) order

```
[24]: import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
series = ts_df['energy_consumption'].dropna()

# Set number of lags (how far back you want to check correlation)
num_lags = 40


plt.figure(figsize=(14, 6))

# ACF plot
plt.subplot(1, 2, 1)
plot_acf(series, lags=num_lags, ax=plt.gca())
plt.title('Autocorrelation (ACF)')

# PACF plot
```

```
plt.subplot(1, 2, 2)
plot_pacf(series, lags=num_lags, ax=plt.gca(), method='ywmm') # 'ywmm' is a
↳ stable method

plt.tight_layout()
plt.show()
```



## 1.5 Interpretation

### 1.5.1 ACF (Left Plot)

- **Lag 1 to ~10:** Gradually decreasing autocorrelation — this is called a **slow decay**.
  - This suggests that the time series is **strongly influenced by its recent past**, and those effects **fade over time**.
  - This kind of pattern is typical in an **AR (AutoRegressive)** process — especially **AR(1)** or **AR(2)**.
- 

### 1.5.2 PACF (Right Plot)

- **Lag 1:** Very strong spike — indicates a **strong direct effect** from lag 1.
- **Lag 2:** Also slightly significant.
- **Lag 3 and beyond:** Mostly within confidence bounds (i.e., not significant).

### 1.5.3 This tells us:

- **Most of the predictive power comes from the first 1 or 2 lags.**
  - After that, past values don't directly contribute much.
-

#### 1.5.4 Conclusion:

Your time series likely follows an **AR(1)** or **AR(2)** model. Here's why:

Indicator	Observation	Suggestion
ACF	Slowly decays	Autoregressive process
PACF	Big spike at lag 1, possibly lag 2	Likely AR(1) or AR(2)